

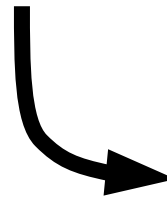
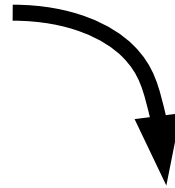


# Machine Learning with Scikit-Learn

Andreas Mueller (NYU Center for Data Science, scikit-learn)

Material: <http://bit.ly/sklsf>

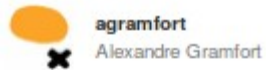
# Me



Classification  
Regression  
Clustering  
Semi-Supervised Learning  
Feature Selection  
Feature Extraction  
Manifold Learning  
Dimensionality Reduction  
Kernel Approximation  
Hyperparameter Optimization  
Evaluation Metrics  
Out-of-core learning

.....





**agramfort**  
Alexandre Gramfort



**AlexanderFabisch**  
Alexander Fabisch



**alextp**  
Alexandre Passos



**amueller**  
Andreas Mueller



**arjoly**  
Arnaud Joly



**bdholt1**  
Brian Holt



**GaelVaroquaux**  
Gael Varoquaux



**glouppe**  
Gilles Louppe



**jakevdp**  
Jake Vanderplas



**jaquesgrobler**  
Jaques Grobler



**jnothman**



**kastnerkyle**  
Kyle Kastner



**bthirion**  
bthirion



**chrisfilo**  
Chris Filo Gorgole...



**cournape**  
David Cournapeau



**duchesnay**  
Duchesnay



**dwf**  
David Warde-Farley



**fabianp**  
Fabian Pedregosa



**kuantkid**  
Wei Li



**larsmans**  
Lars



**lucidfrontier45**  
Shiqiao Du



**mblondel**  
Mathieu Blondel



**MechCoder**  
Manoj Kumar



**ndawe**  
Noel Dawe



**NelleV**  
Varoquaux



**ogrisel**  
Olivier Grisel



**paolo-losi**  
Paolo Losi



**pprett**  
Peter Prettenhofer



**robertlayton**  
Robert Layton



**ronw**  
Ron Weiss



**satra**  
Satrajit Ghosh



**sklearn-ci**



**vene**  
Vlad Niculae



**VirgileFritsch**  
Virgile Fritsch



**vmichel**  
Vincent Michel



**yarikoptic**  
Yaroslav Halchenko



# Get the notebooks!

nbviewer   FAQ   IPython



pydata-nyc-advanced-sklearn / Chapter 0 - Reminder.ipynb /

## Scikit-Learn is simple

### Classification

```
In [4]: from sklearn.datasets import load_iris
        from sklearn.cross_validation import train_test_split

iris = load_iris()
X, y = iris.data, iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
In [5]: from sklearn.svm import SVC
        clf = SVC()
        clf.fit(X_train, y_train)
        y_pred = clf.predict(X_test)
```

### Transformations

```
In [6]: from sklearn.decomposition import PCA
```

```
In [7]: pca = PCA(n_components=2)
        pca.fit(X)
        X_pca = pca.transform(X)
```

## Tools

### Cross-validation scoring

```
In [30]: from sklearn.cross_validation import cross_val_score, StratifiedKFold
        scores = cross_val_score(SVC(), X_train, y_train, cv=5)
        print(scores)

[ 0.95652174  1.          0.95652174  0.91304348  0.9         ]
```

<http://bit.ly/sklsf>

## Documentation of scikit-learn 0.17

### Quick Start

A very short introduction into machine learning problems and how to solve them using scikit-learn. Introduced basic concepts and conventions.

### User Guide

The main documentation. This contains an in-depth description of all algorithms and how to apply them.

### Other Versions

- [scikit-learn 0.18 \(development\)](#)
- [scikit-learn 0.17 \(stable\)](#)
- [scikit-learn 0.16](#)
- [scikit-learn 0.15](#)

### Tutorials

Useful tutorials for developing a feel for some of scikit-learn's applications in the machine learning field.

### API

The exact API of all functions and classes, as given by the docstrings. The API documents expected types and allowed features for all functions, and all parameters available for the algorithms.

### Additional Resources

Talks given, slide-sets and other information relevant to scikit-learn.

### Contributing

Information on how to contribute. This also contains useful information for advanced users, for example how to build their own estimators.

### Flow Chart

A graphical overview of basic areas of machine learning, and guidance which kind of algorithms to use in a given situation.

### FAQ

Frequently asked questions about the project and contributing.

<http://scikit-learn.org/>

Hi Andy,

I just received an email from the first tutorial speaker, presenting right before you, saying he's ill and won't be able to make it.

I know you have already committed yourself to two presentations, but is there anyway you could increase your tutorial time slot, maybe just offer time to try out what you've taught? Otherwise I have to do some kind of modern dance interpretation of Python in data :-)  
-Leah

Hi Andreas,

I am very interested in your Machine Learning background. I work for X Recruiting who have been engaged by Z, a worldwide leading supplier of Y. We are expanding the core engineering team and we are looking for really passionate engineers who want to create their own story and help millions of people.

Can we find a time for a call to chat for a few minutes about this?

Thanks

Hi Andy,

I just received an email from the first tutorial speaker, presenting me, saying he's ill and won't be

I know you have a lot to offer yourself to two presentations. I think you could increase your audience, maybe just offer time to try out what you've taught? Otherwise I have to do some kind of modern dance interpretation of Python in data :-)

-Leah



Hi Andreas,

I am very interested in your Machine Learning background. I work for X Recruiting who have been engaged by Z, a worldwide leading supplier of Y. We are expanding the core engineering team and we are looking for really passionate engineers who can tell their own story and help millions of

Can we find a time for a few minutes about this?

Thanks





# Doing Machine Learning With Scikit-Learn

# Representing Data

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

# Representing Data

one sample

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

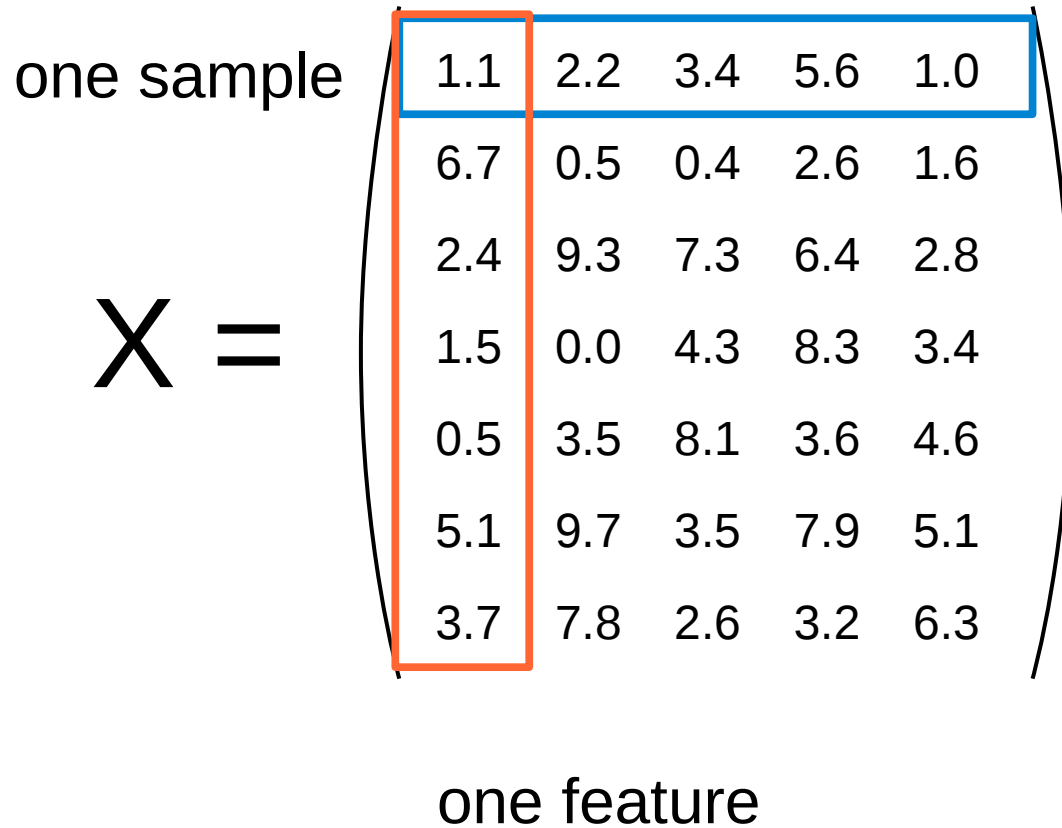
# Representing Data

one sample

$X =$

|     |     |     |     |     |
|-----|-----|-----|-----|-----|
| 1.1 | 2.2 | 3.4 | 5.6 | 1.0 |
| 6.7 | 0.5 | 0.4 | 2.6 | 1.6 |
| 2.4 | 9.3 | 7.3 | 6.4 | 2.8 |
| 1.5 | 0.0 | 4.3 | 8.3 | 3.4 |
| 0.5 | 3.5 | 8.1 | 3.6 | 4.6 |
| 5.1 | 9.7 | 3.5 | 7.9 | 5.1 |
| 3.7 | 7.8 | 2.6 | 3.2 | 6.3 |

one feature



# Representing Data

one sample

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

one feature

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

outputs / labels

# Training and Testing Data

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix} \quad y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

# Training and Testing Data

training set

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

test set

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

# Training and Testing Data

training set

|       |     |     |     |     |     |
|-------|-----|-----|-----|-----|-----|
| $X =$ | 1.1 | 2.2 | 3.4 | 5.6 | 1.0 |
|       | 6.7 | 0.5 | 0.4 | 2.6 | 1.6 |
|       | 2.4 | 9.3 | 7.3 | 6.4 | 2.8 |
|       | 1.5 | 0.0 | 4.3 | 8.3 | 3.4 |
|       | 0.5 | 3.5 | 8.1 | 3.6 | 4.6 |
|       | 5.1 | 9.7 | 3.5 | 7.9 | 5.1 |
|       | 3.7 | 7.8 | 2.6 | 3.2 | 6.3 |
|       |     |     |     |     |     |
|       |     |     |     |     |     |
|       |     |     |     |     |     |

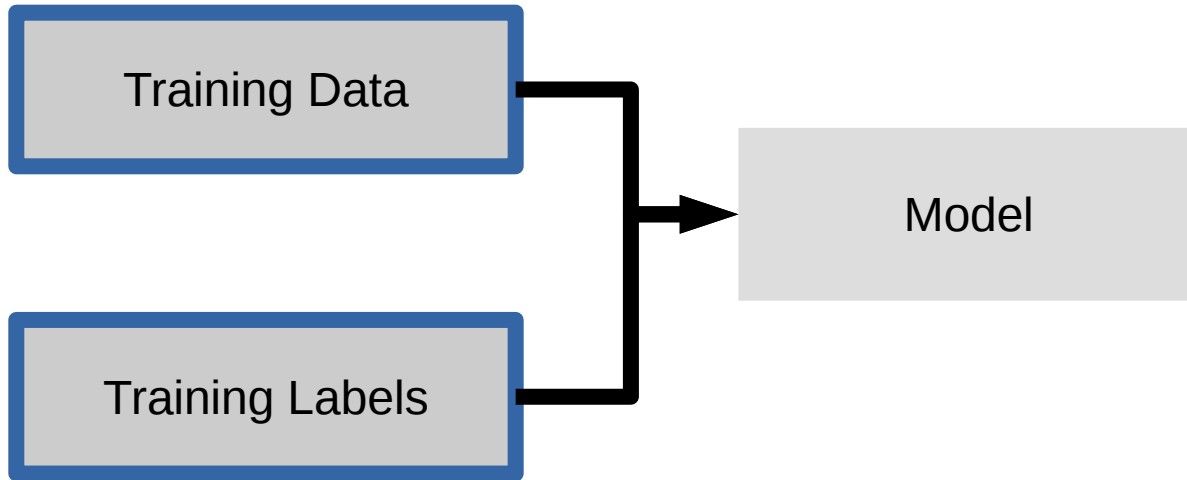
test set

|       |     |
|-------|-----|
| $y =$ | 1.6 |
|       | 2.7 |
|       | 4.4 |
|       | 0.5 |
|       | 0.2 |
|       | 5.6 |
|       | 6.7 |
|       |     |

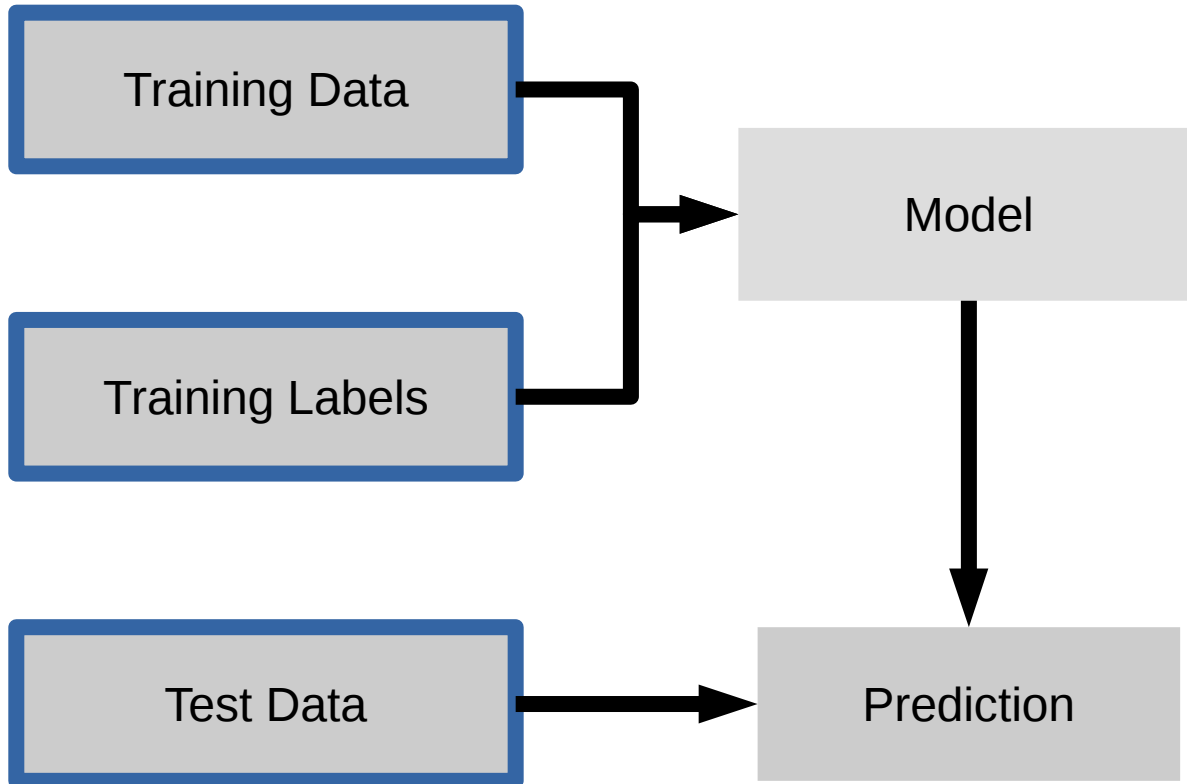
```
from sklearn.cross_validation import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y)
```



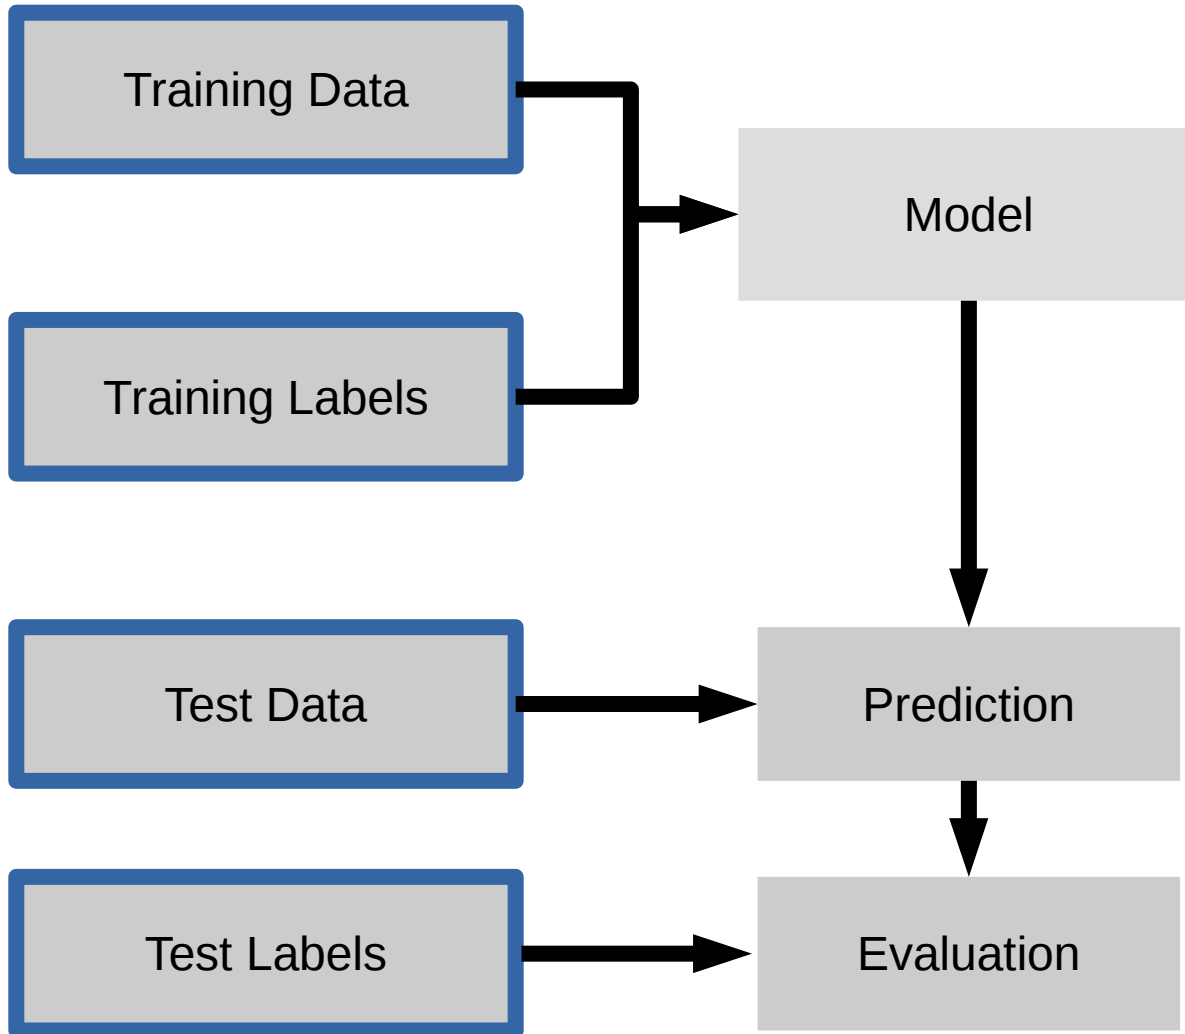
# Supervised Machine Learning



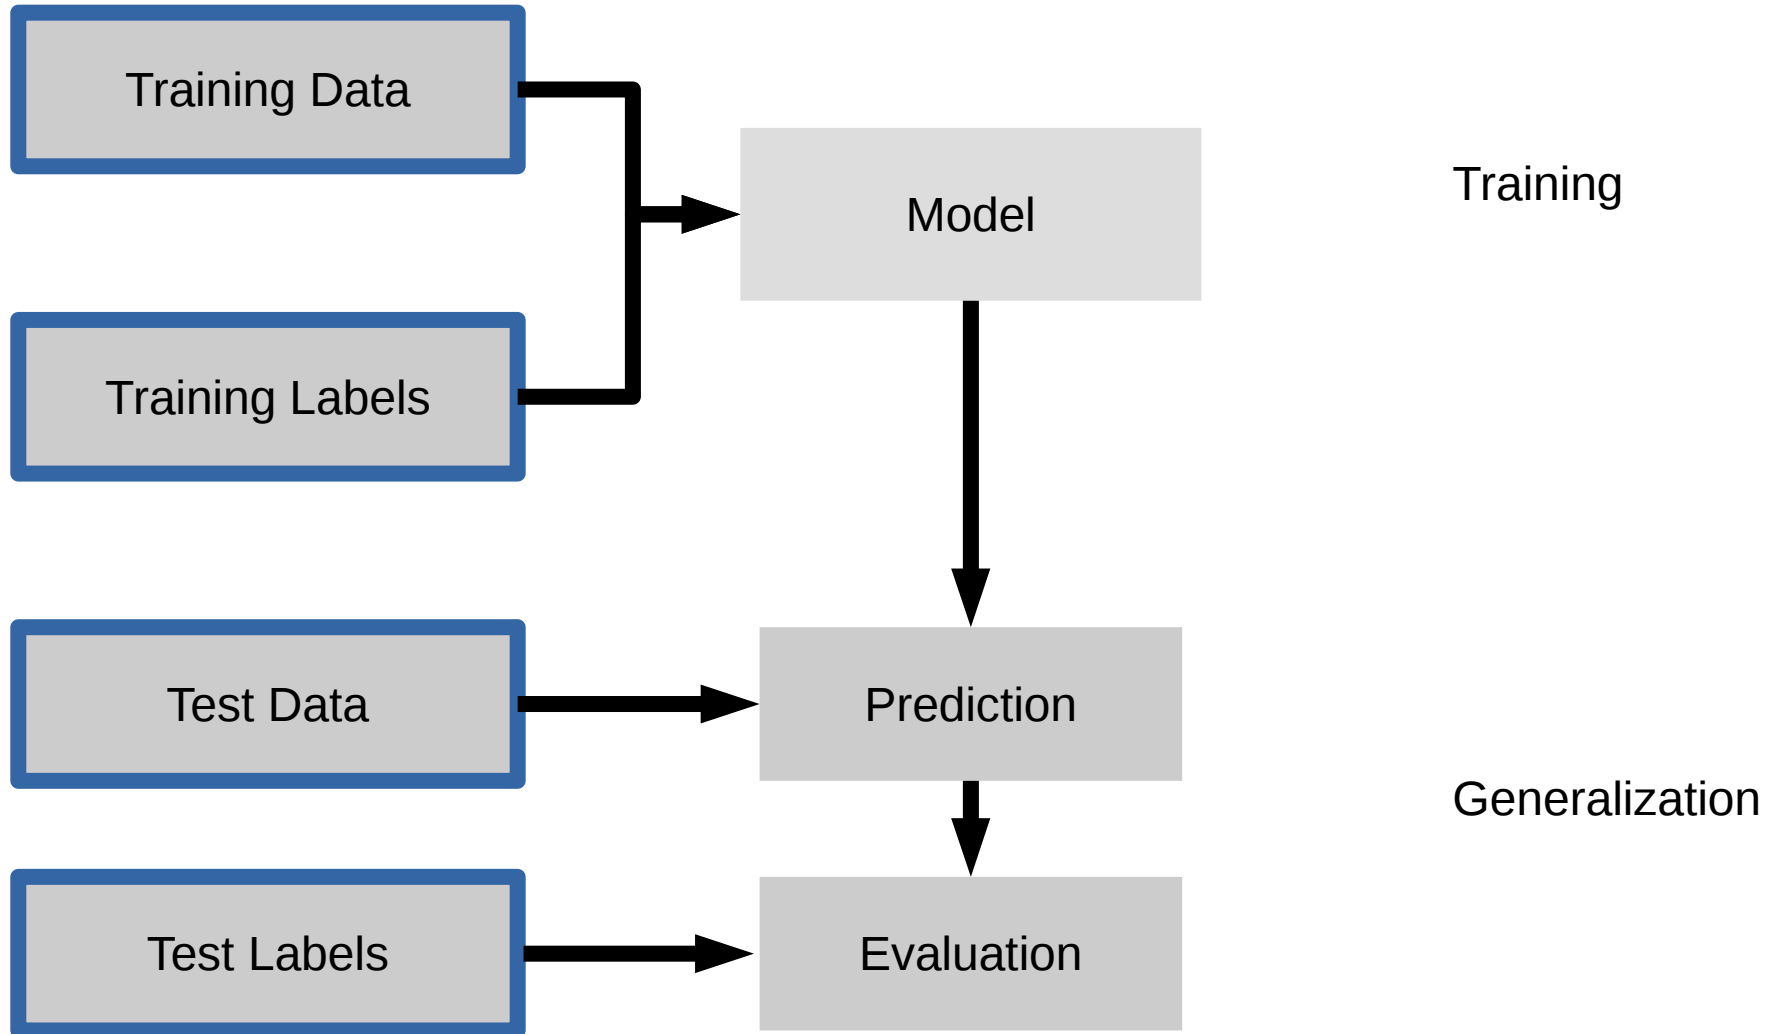
# Supervised Machine Learning



# Supervised Machine Learning

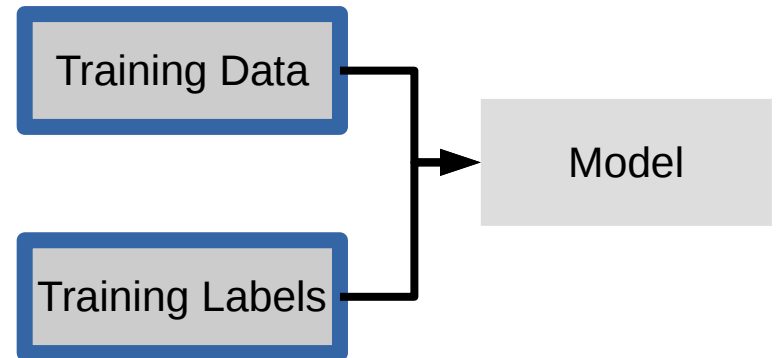


# Supervised Machine Learning



```
clf = RandomForestClassifier()
```

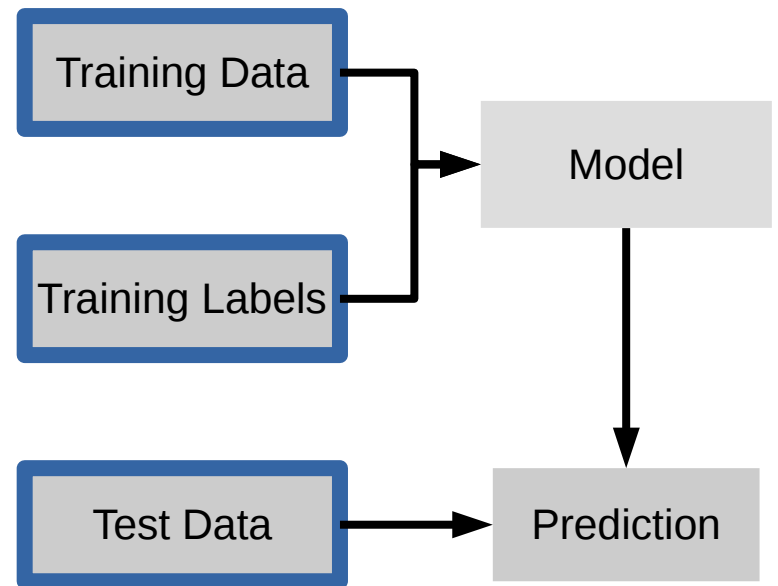
```
clf.fit(X_train, y_train)
```



```
clf = RandomForestClassifier()
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

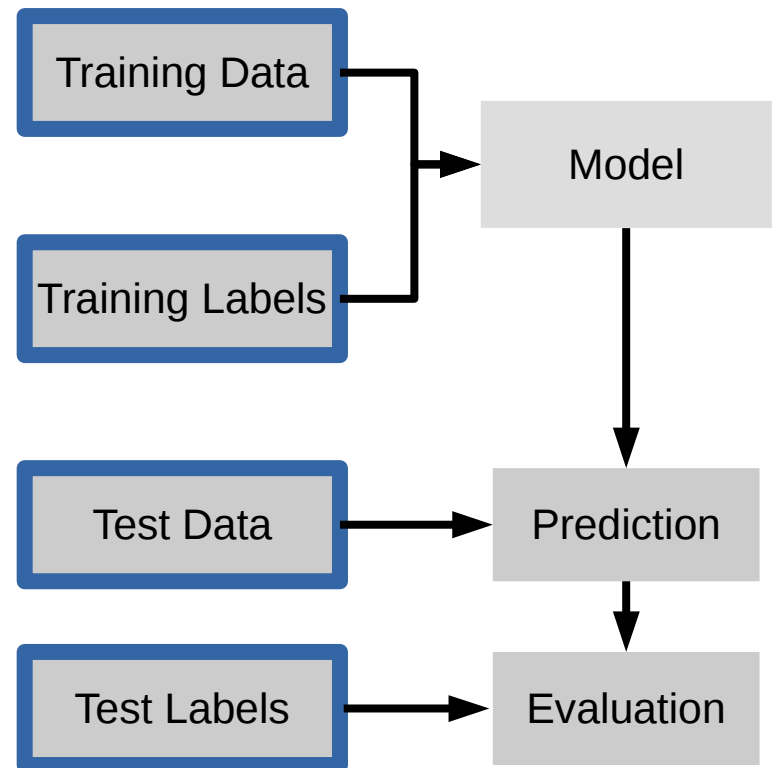


```
clf = RandomForestClassifier()
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
clf.score(X_test, y_test)
```

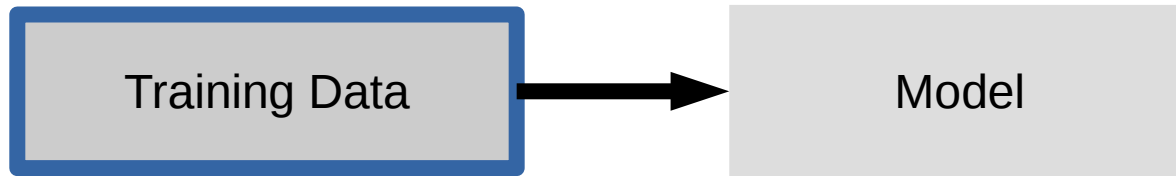


# IPython Notebook:

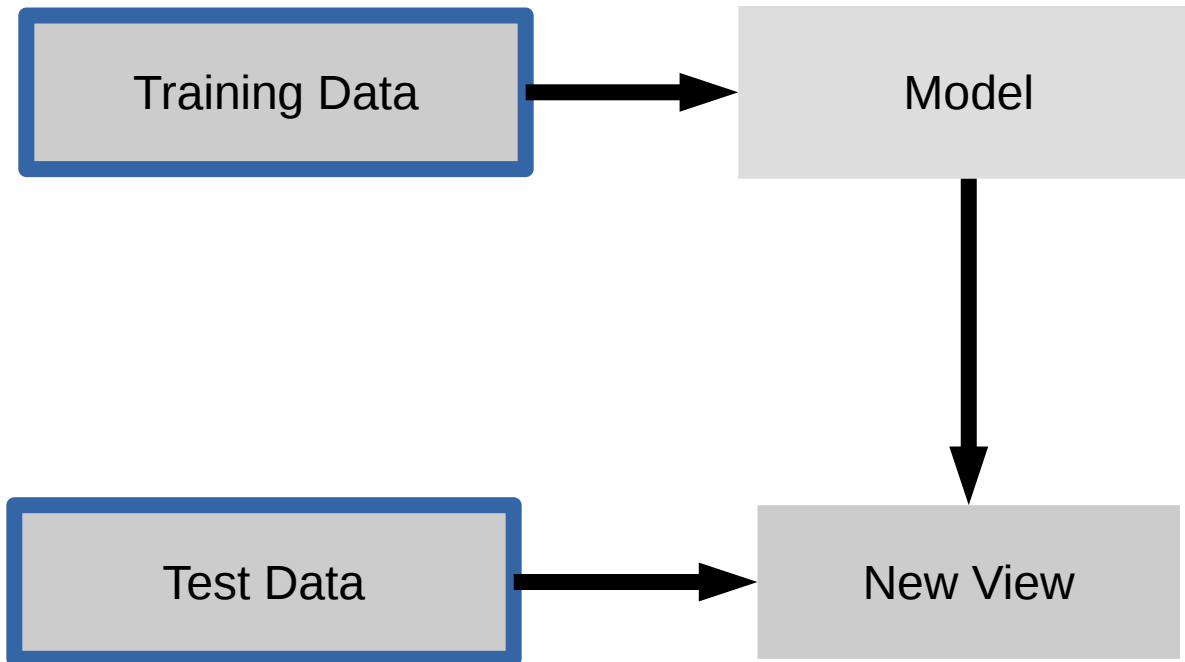
## Part 1 - Introduction to Scikit-learn



# Unsupervised Machine Learning



# Unsupervised Machine Learning

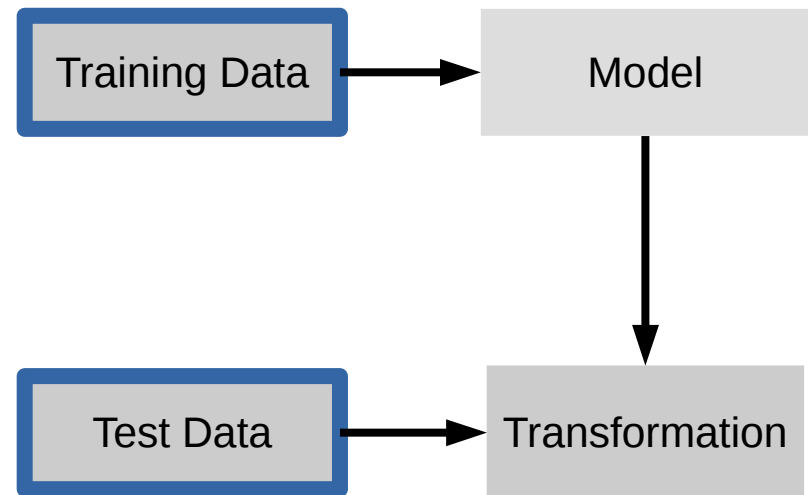


# Unsupervised Transformations

```
pca = PCA()
```

```
pca.fit(X_train)
```

```
X_new = pca.transform(X_test)
```



# IPython Notebook:

## Part 2 – Unsupervised Transformers

# Basic API

**`estimator.fit(X, [y])`**

**`estimator.predict`**

**`estimator.transform`**

---

Classification

Preprocessing

Regression

Dimensionality reduction

Clustering

Feature selection

Feature extraction

All Data

Training data

Test data

All Data

Training data

Test data

Fold 1

Fold 2

Fold 3

Fold 4

Fold 5

All Data

Training data      Test data

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 1

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5



All Data

Training data      Test data

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 1

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 2

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

All Data

Training data      Test data

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 1

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 2

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 3

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 4

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

Split 5

Fold 1      Fold 2      Fold 3      Fold 4      Fold 5

# IPython Notebook: Part 3 - Cross-validation

```
In [2]: clf = SVC()  
        clf.fit(X_train, y_train)  
        y_pred = clf.predict(X_test)
```

```
In [2]: clf = SVC()  
        clf.fit(X_train, y_train)
```

```
SVC(self, C=1.0, kernel='rbf', degree=3, gamma=0.0, coef0=0.0,  
     shrinking=True, probability=False, tol=0.001, cache_size=200,  
     class_weight=None, verbose=False, max_iter=-1, random_state=None)
```

All Data

Training data

Test data

All Data

Training data      Test data

Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 1   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 2   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 3   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 4   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Split 5   Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

Test data

All Data

Training data      Test data

Fold 1   Fold 2   Fold 3   Fold 4   Fold 5

|         |        |        |        |        |        |
|---------|--------|--------|--------|--------|--------|
| Split 1 | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 |
| Split 2 | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 |
| Split 3 | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 |
| Split 4 | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 |
| Split 5 | Fold 1 | Fold 2 | Fold 3 | Fold 4 | Fold 5 |

Finding Parameters

Final evaluation

Test data



```
SVC(C=0.001,  
gamma=0.001)
```

SVC(C=0.001,  
gamma=0.001)

SVC(C=0.01,  
gamma=0.001)

SVC(C=0.1,  
gamma=0.001)

SVC(C=1,  
gamma=0.001)


SVC(C=10,  
gamma=0.001)

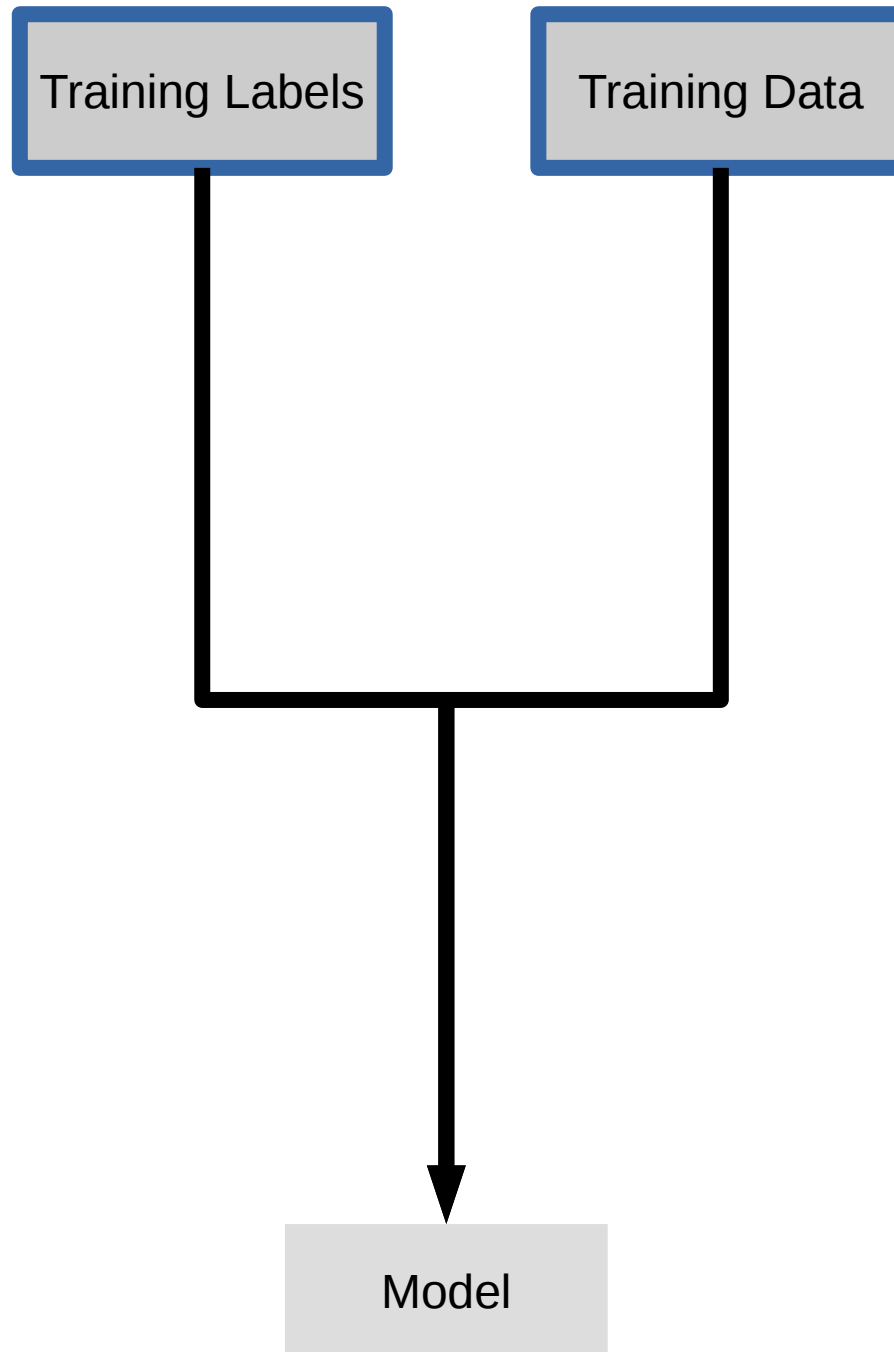
|                              |                             |                            |                          |                           |
|------------------------------|-----------------------------|----------------------------|--------------------------|---------------------------|
| SVC(C=0.001,<br>gamma=0.001) | SVC(C=0.01,<br>gamma=0.001) | SVC(C=0.1,<br>gamma=0.001) | SVC(C=1,<br>gamma=0.001) | SVC(C=10,<br>gamma=0.001) |
| SVC(C=0.001,<br>gamma=0.01)  | SVC(C=0.01,<br>gamma=0.01)  | SVC(C=0.1,<br>gamma=0.01)  | SVC(C=1,<br>gamma=0.01)  | SVC(C=10,<br>gamma=0.01)  |

|                              |                             |                            |                          |                           |
|------------------------------|-----------------------------|----------------------------|--------------------------|---------------------------|
| SVC(C=0.001,<br>gamma=0.001) | SVC(C=0.01,<br>gamma=0.001) | SVC(C=0.1,<br>gamma=0.001) | SVC(C=1,<br>gamma=0.001) | SVC(C=10,<br>gamma=0.001) |
| SVC(C=0.001,<br>gamma=0.01)  | SVC(C=0.01,<br>gamma=0.01)  | SVC(C=0.1,<br>gamma=0.01)  | SVC(C=1,<br>gamma=0.01)  | SVC(C=10,<br>gamma=0.01)  |
| SVC(C=0.001,<br>gamma=0.1)   | SVC(C=0.01,<br>gamma=0.1)   | SVC(C=0.1,<br>gamma=0.1)   | SVC(C=1,<br>gamma=0.1)   | SVC(C=10,<br>gamma=0.1)   |


|                              |                             |                            |                          |                           |
|------------------------------|-----------------------------|----------------------------|--------------------------|---------------------------|
| SVC(C=0.001,<br>gamma=0.001) | SVC(C=0.01,<br>gamma=0.001) | SVC(C=0.1,<br>gamma=0.001) | SVC(C=1,<br>gamma=0.001) | SVC(C=10,<br>gamma=0.001) |
| SVC(C=0.001,<br>gamma=0.01)  | SVC(C=0.01,<br>gamma=0.01)  | SVC(C=0.1,<br>gamma=0.01)  | SVC(C=1,<br>gamma=0.01)  | SVC(C=10,<br>gamma=0.01)  |
| SVC(C=0.001,<br>gamma=0.1)   | SVC(C=0.01,<br>gamma=0.1)   | SVC(C=0.1,<br>gamma=0.1)   | SVC(C=1,<br>gamma=0.1)   | SVC(C=10,<br>gamma=0.1)   |
| SVC(C=0.001,<br>gamma=1)     | SVC(C=0.01,<br>gamma=1)     | SVC(C=0.1,<br>gamma=1)     | SVC(C=1,<br>gamma=1)     | SVC(C=10,<br>gamma=1)     |
| SVC(C=0.001,<br>gamma=10)    | SVC(C=0.01,<br>gamma=10)    | SVC(C=0.1,<br>gamma=10)    | SVC(C=1,<br>gamma=10)    | SVC(C=10,<br>gamma=10)    |


# IPython Notebook: Part 4 – Grid Searches

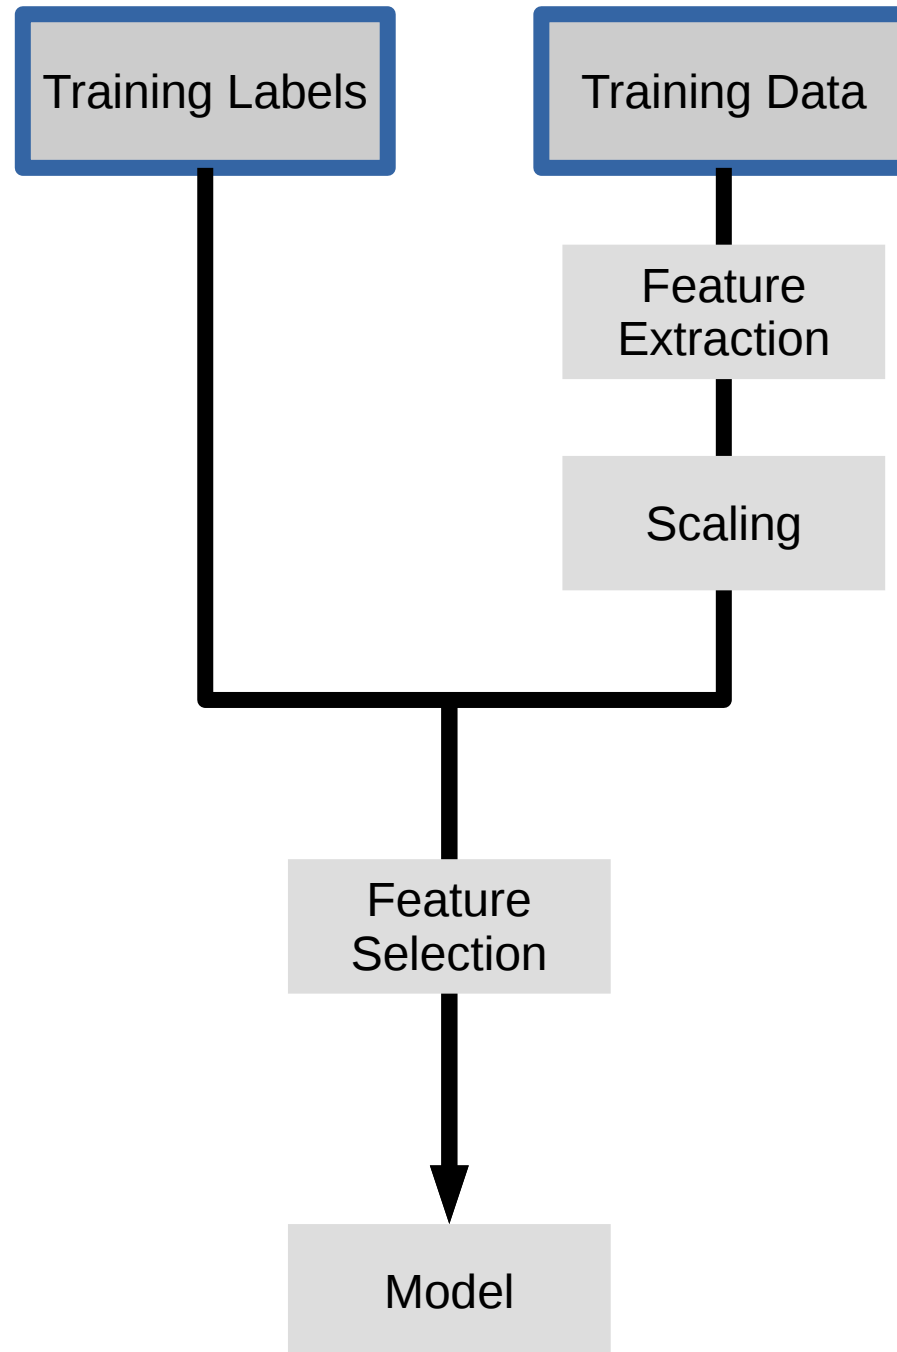


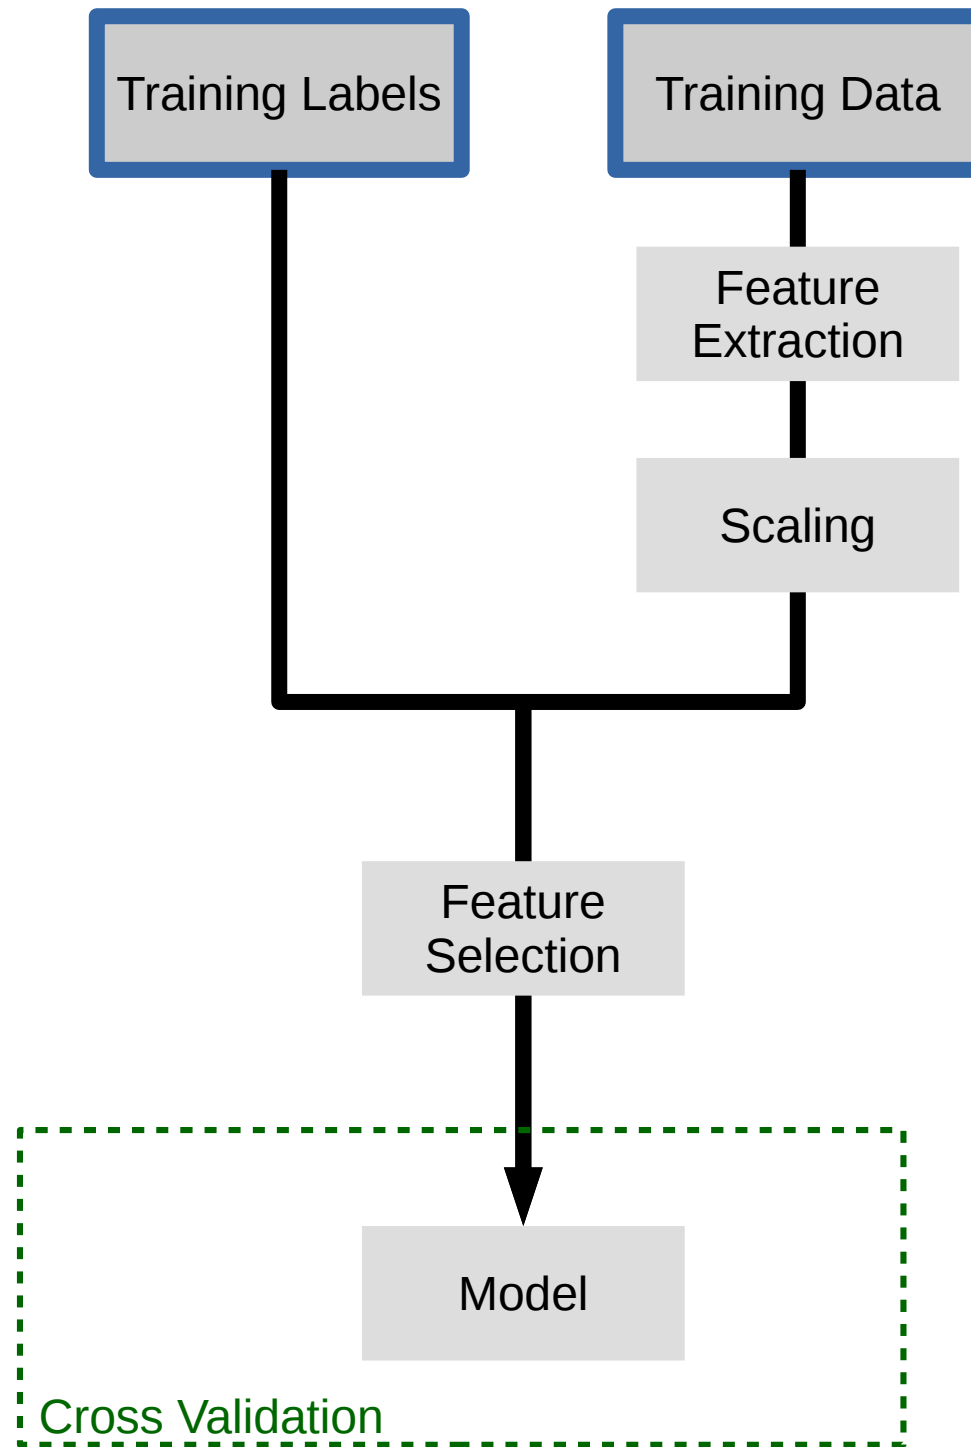


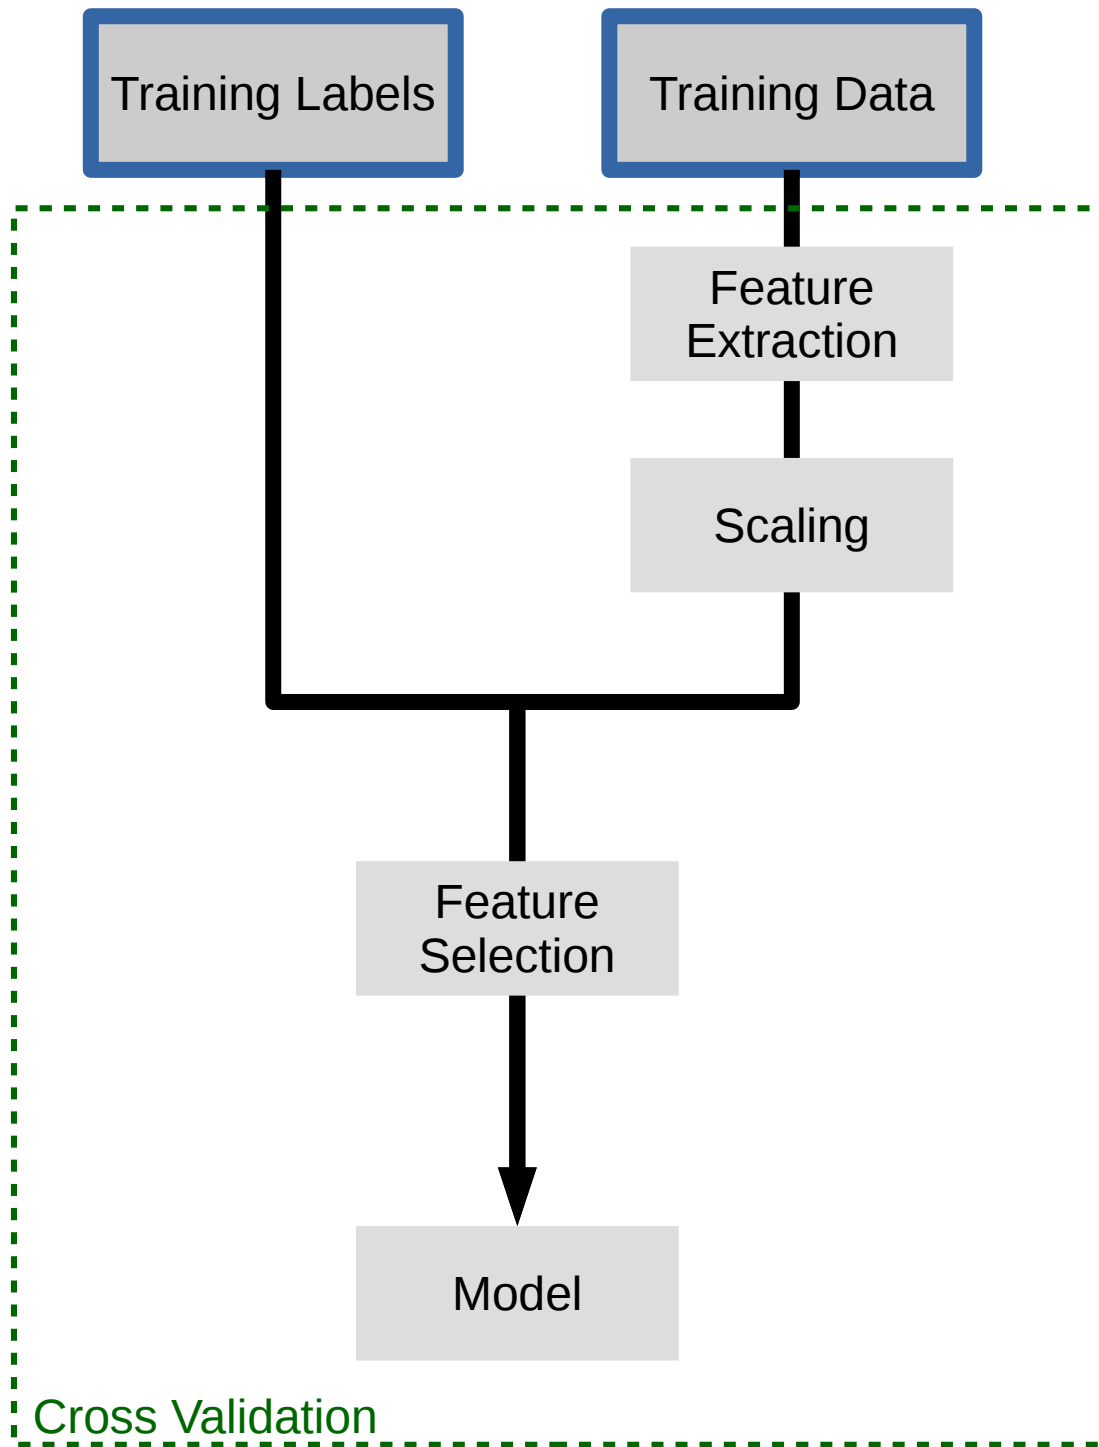












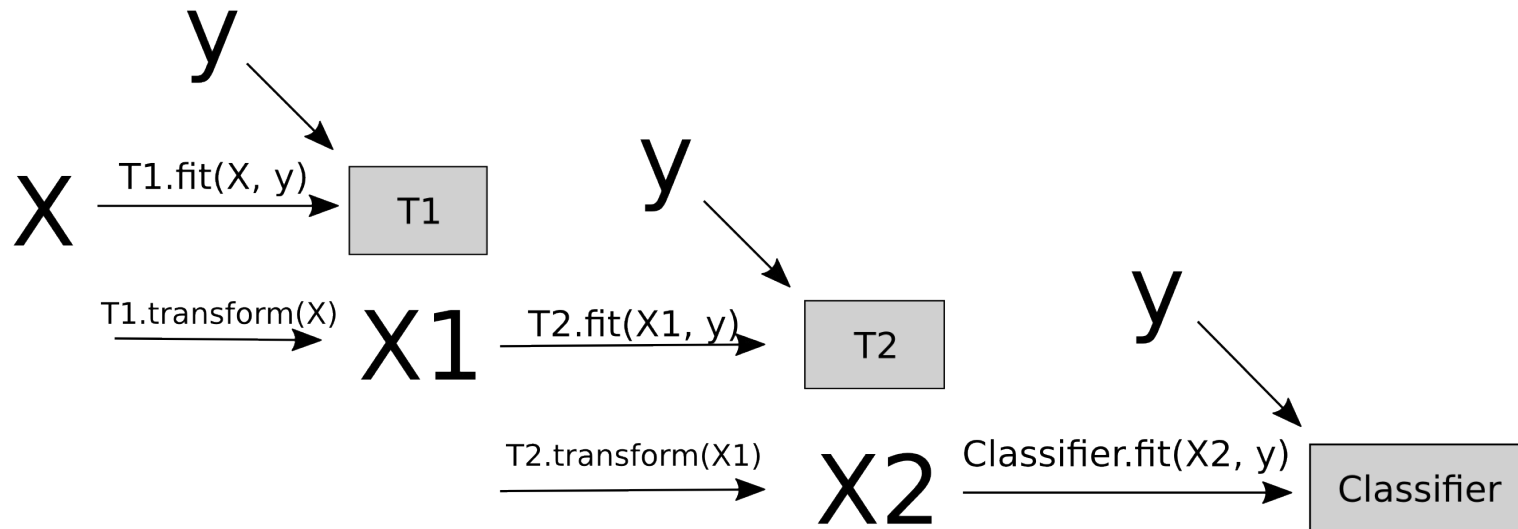
# Pipelines

# Pipelines

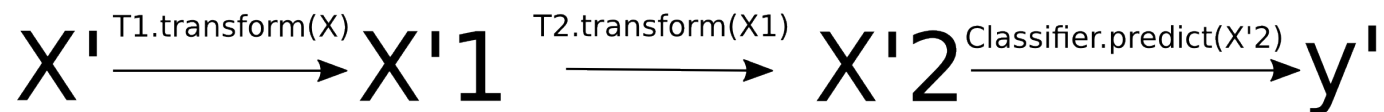
```
pipe = make_pipeline(T1(), T2(), Classifier())
```



```
pipe.fit(X, y)
```



```
pipe.predict(X)
```



# IPython Notebook:

## Part 5 - Preprocessing and Pipelines



Do cross-validation over all steps jointly.  
Keep a separate test set until the very end.

## Sample application: Sentiment Analysis

# IMDB Movie Reviews Data

## **Review:**

One of the worst movies I've ever rented. Sorry it had one of my favorite actors on it (Travolta) in a nonsense role. In fact, anything made sense in this movie.

Who can say there was true love between Eddy and Maureen?  
Don't you remember the beginning of the movie ?

Is she so lovely? Ask her daughters. I don't think so.

**Label:** negative

**Training data:** 12500 positive, 12500 negative

# Bag Of Word Representations

`CountVectorizer / TfidfVectorizer`

# Bag Of Word Representations

`CountVectorizer / TfidfVectorizer`

`"This is how you get ants."`

# Bag Of Word Representations

CountVectorizer / TfidfVectorizer

"This is how you get ants."

tokenizer

↓  
['this', 'is', 'how', 'you', 'get', 'ants']

# Bag Of Word Representations

CountVectorizer / TfidfVectorizer

"This is how you get ants."

tokenizer

['this', 'is', 'how', 'you', 'get', 'ants']

Build a vocabulary over all documents

['aardvak', 'amsterdam', 'ants', ... 'you', 'your', 'zyxst']

# Bag Of Word Representations

CountVectorizer / TfidfVectorizer

"This is how you get ants."

tokenizer

['this', 'is', 'how', 'you', 'get', 'ants']

Build a vocabulary over all documents

['aardvak', 'amsterdam', 'ants', ... 'you', 'your', 'zyxst']

Sparse matrix encoding

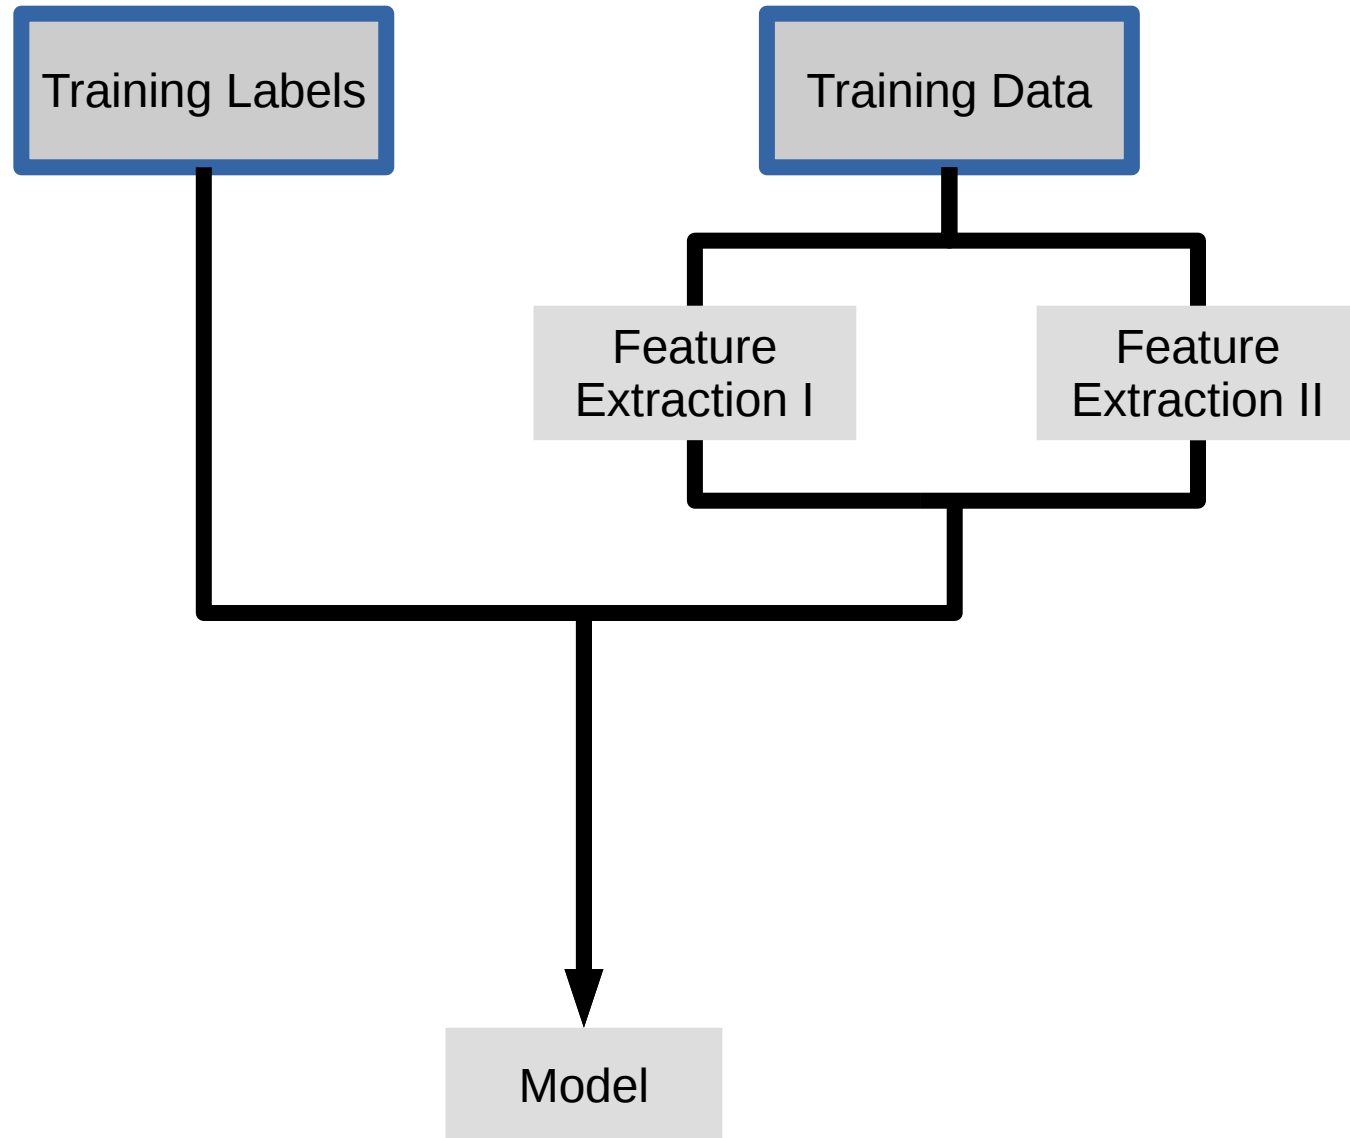
| aardvak   | ants | get | you | zyxst |
|---|------|-----|-----|-------|
| [0, ..., 0, 1, 0, ... , 0, 1 , 0, ..., 0, 1, 0, ..... | 0]   |     |     |       |



# IPython Notebook:


## Part 6 - Working With Text Data

# Feature Union




# IPython Notebook: Part 7 – FeatureUnion


# Overfitting and Underfitting



# Overfitting and Underfitting



# Overfitting and Underfitting






# Linear SVM

$$\hat{y} = \text{sign}(w_0 + \sum_i w_i x_i)$$

# Linear SVM

$$\hat{y} = \text{sign}(w_0 + \sum_i w_i x_i)$$





# (RBF) Kernel SVM


$$\hat{y} = \text{sign}(\alpha_0 + \sum_j \alpha_j y_j k(\mathbf{x}^{(j)}, \mathbf{x}))$$

# (RBF) Kernel SVM

$$\hat{y} = \text{sign}(\alpha_0 + \sum_j \alpha_j y_j k(\mathbf{x}^{(j)}, \mathbf{x}))$$
$$k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$$

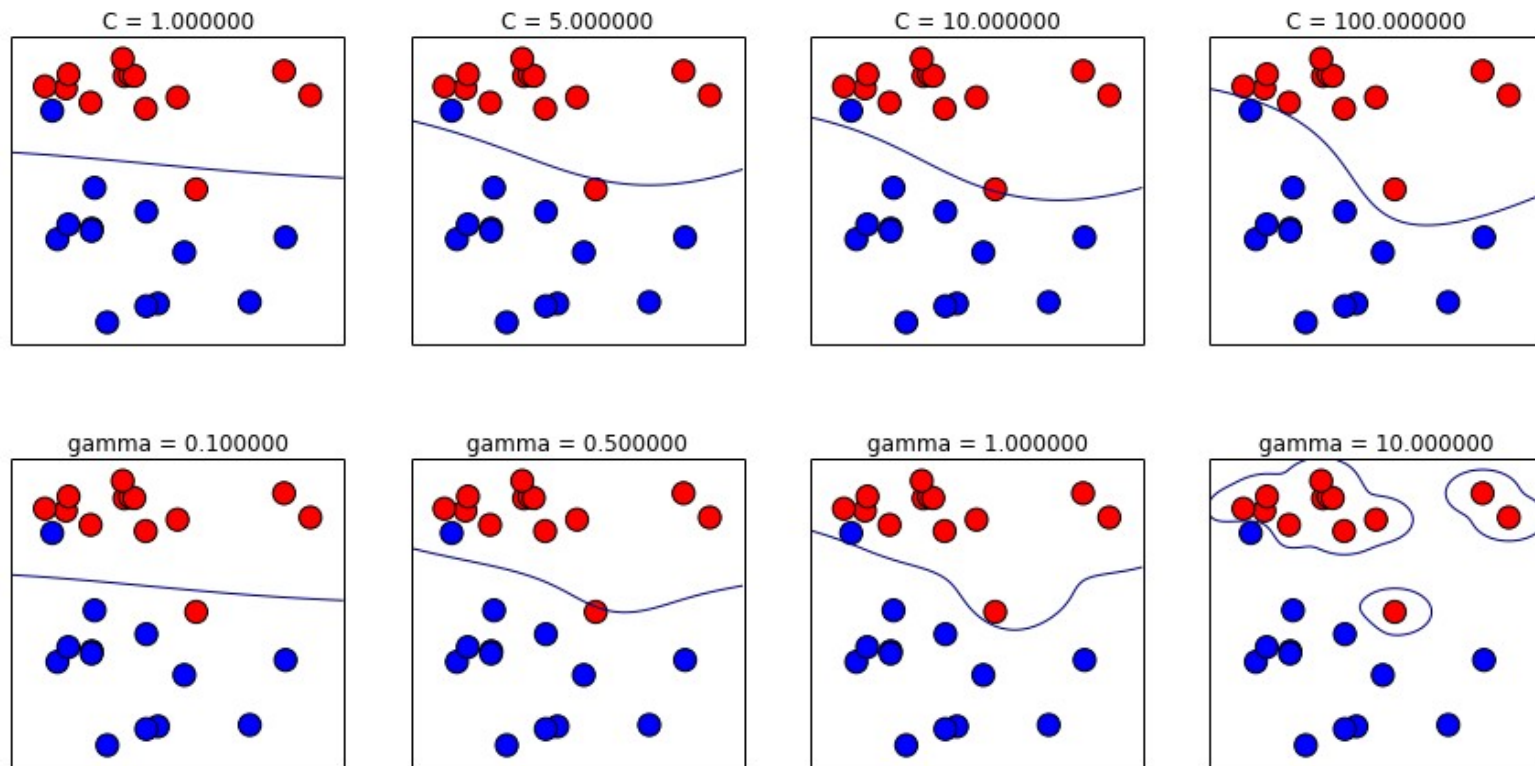
# (RBF) Kernel SVM

$$\hat{y} = \text{sign}(\alpha_0 + \sum_j \alpha_j y_j k(\mathbf{x}^{(j)}, \mathbf{x}))$$
$$k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$$




# (RBF) Kernel SVM


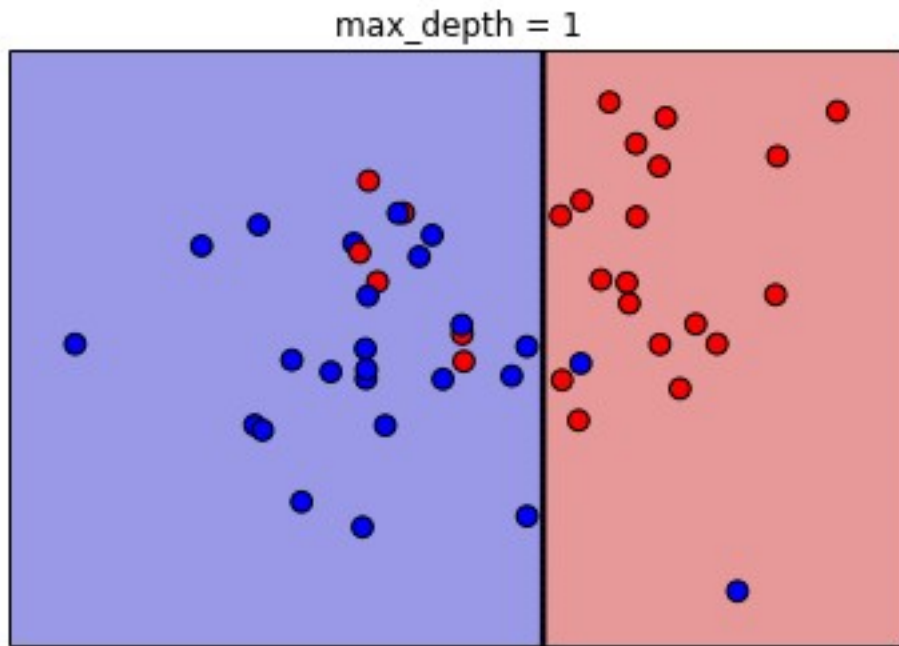
$$\hat{y} = \text{sign}(\alpha_0 + \sum_j \alpha_j y_j k(\mathbf{x}^{(j)}, \mathbf{x}))$$
$$k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|^2)$$





# Decision Trees



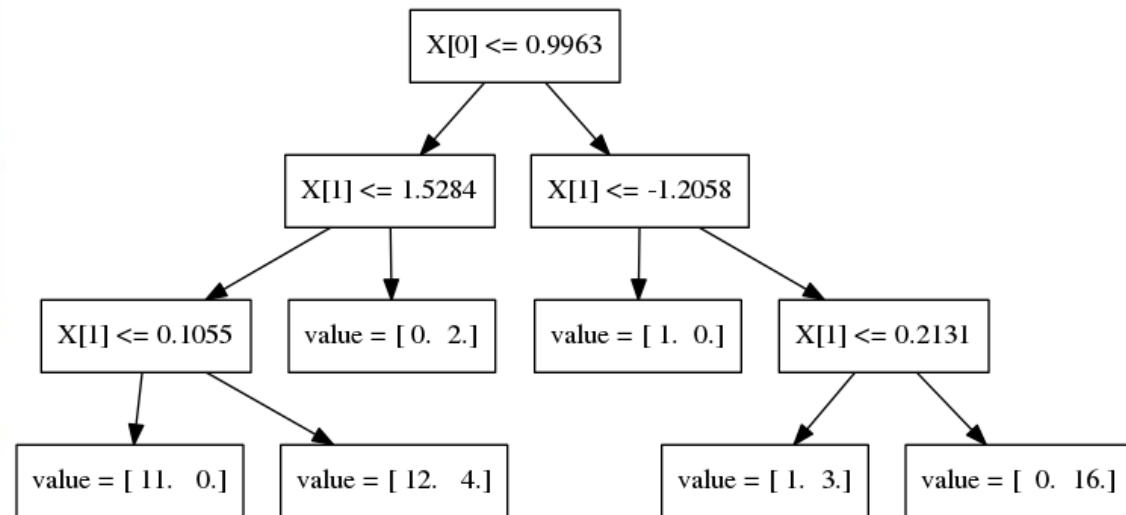

# Decision Trees



# Decision Trees




# Decision Trees

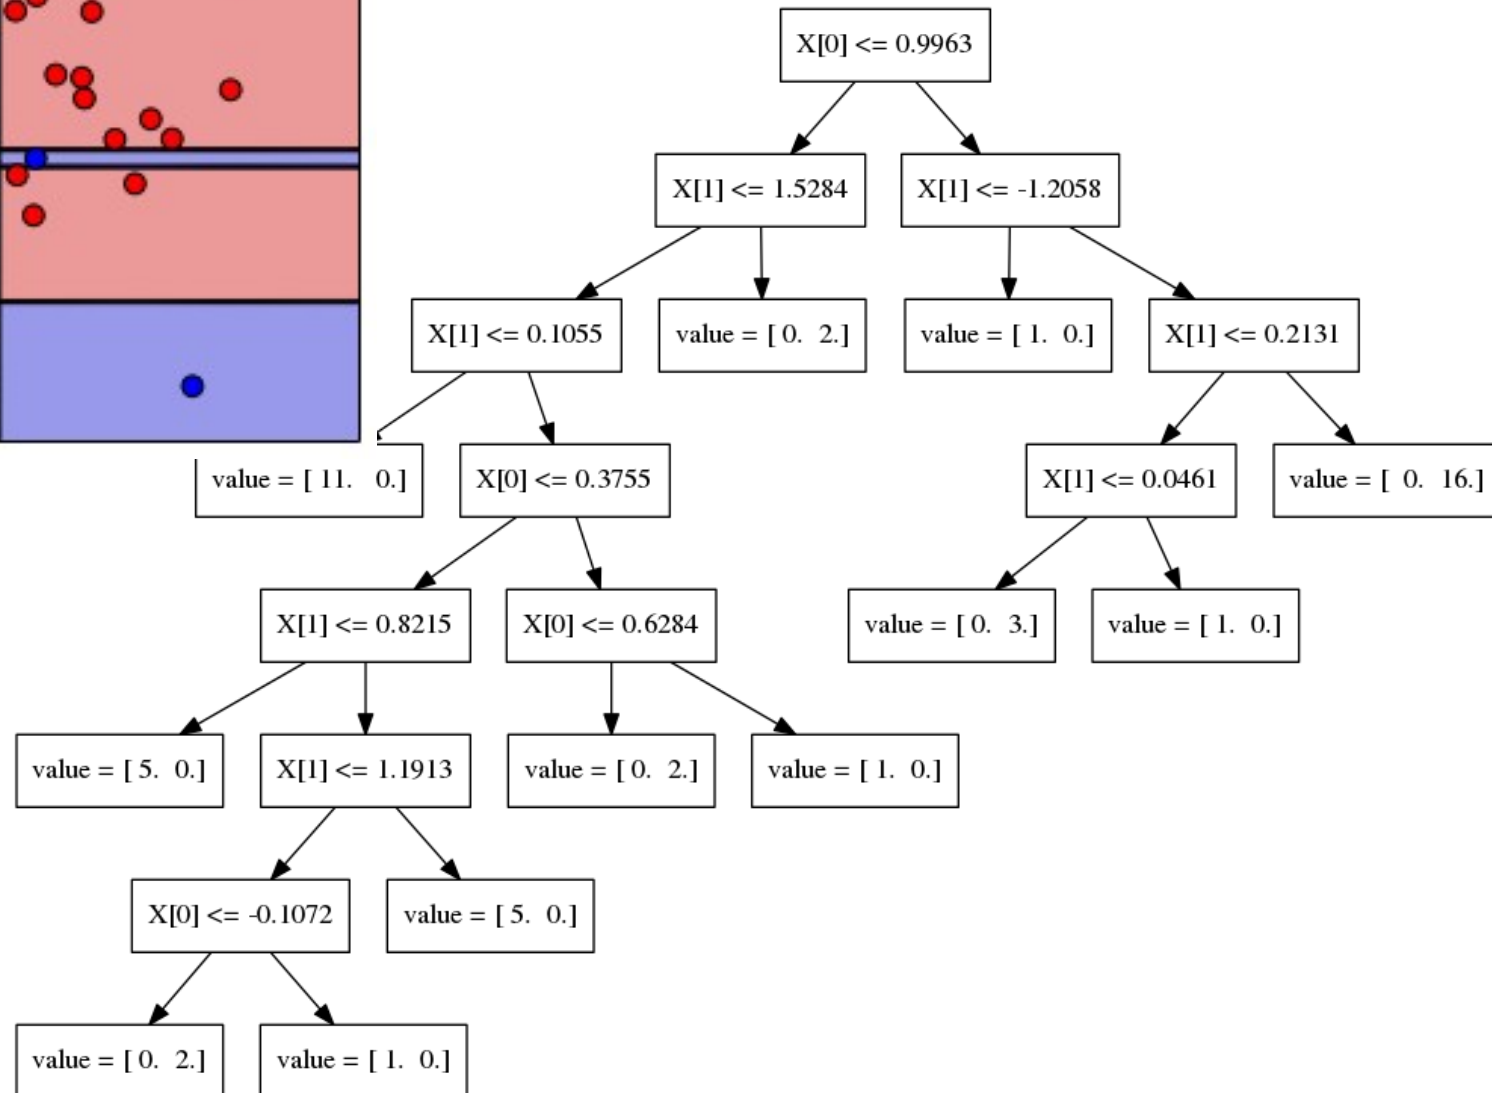






# Decision Trees






# Decision Trees







# Random Forests



# Random Forests




# Random Forests



# Validation Curves

```
train_scores, test_scores = validation_curve(SVC(), X, y,  
      param_name="gamma", param_range=param_range)
```



# Scaling Up

# Three regimes of data

- Fits in RAM
- Fits on a Hard Drive
- Doesn't fit on a single PC



# Three regimes of data

- Fits in RAM (up to 256 GB?)
- Fits on a Hard Drive (up to 6TB?)
- Doesn't fit on a single PC

# Nobody ever got fired for using Hadoop on a cluster

Antony Rowstron, Dushyanth Narayanan, Austin Donnelly, Greg O'Shea, and Andrew Douglas

10 April 2012


|   | vCPU | ECU | Memory (GiB) | Instance Storage (GB) | Linux/UNIX Usage |
|---|------|-----|--------------|-----------------------|------------------|
| <b>Memory Optimized - Current Generation</b>  |      |     |              |                       |                  |
| r3.large                                      | 2    | 6.5 | 15           | 1 x 32 SSD            | \$0.195 per Hour |
| r3.xlarge                                     | 4    | 13  | 30.5         | 1 x 80 SSD            | \$0.39 per Hour  |
| r3.2xlarge                                    | 8    | 26  | 61           | 1 x 160 SSD           | \$0.78 per Hour  |
| r3.4xlarge                                    | 16   | 52  | 122          | 1 x 320 SSD           | \$1.56 per Hour  |
| r3.8xlarge                                    | 32   | 104 | 244          | 2 x 320 SSD           | \$3.12 per Hour  |
| <b>Storage Optimized - Current Generation</b> |      |     |              |                       |                  |
| i2.xlarge                                     | 4    | 14  | 30.5         | 1 x 800 SSD           | \$0.938 per Hour |
| i2.2xlarge                                    | 8    | 27  | 61           | 2 x 800 SSD           | \$1.876 per Hour |
| i2.4xlarge                                    | 16   | 53  | 122          | 4 x 800 SSD           | \$3.751 per Hour |
| i2.8xlarge                                    | 32   | 104 | 244          | 8 x 800 SSD           | \$7.502 per Hour |

"256Gb ought to be enough for anybody."  
- me

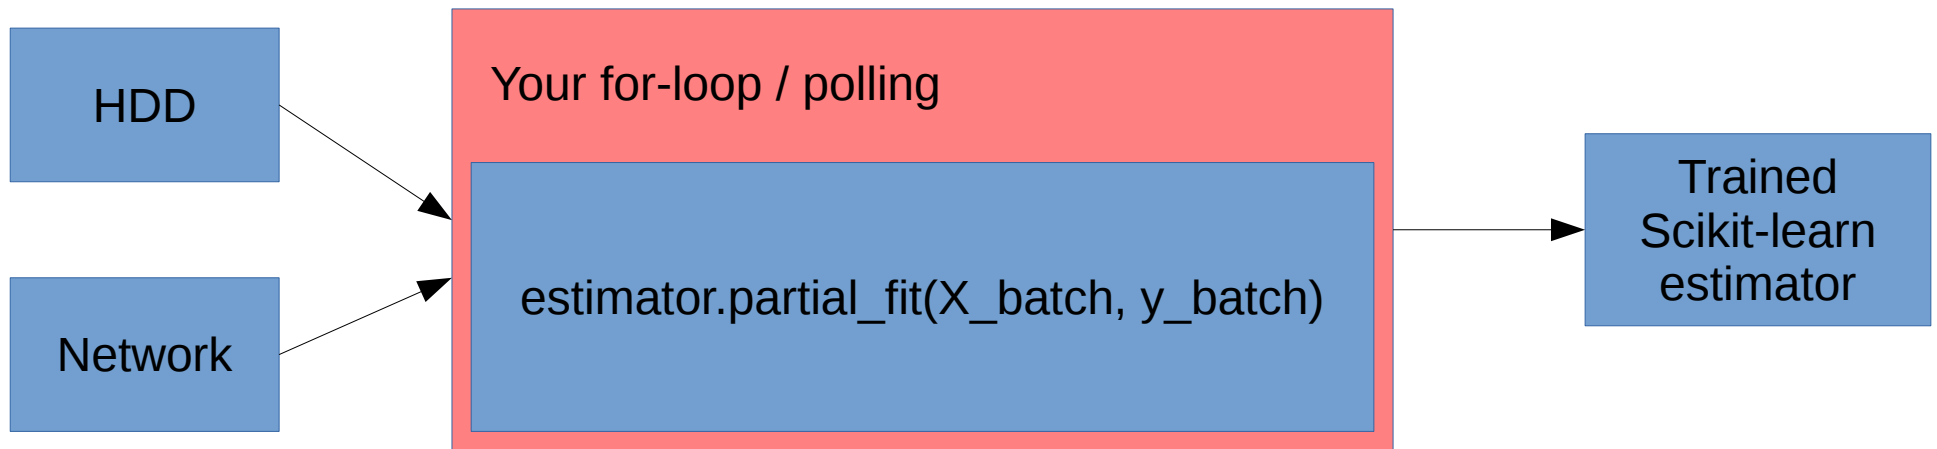
"256Gb ought to be enough for anybody."  
- me

(for machine learning)

# Subsample!



# The scikit-learn way





# Supported Algorithms

- All `SGDClassifier` derivatives
- Naive Bayes
- `MinibatchKMeans`
- Birch
- `IncrementalPCA`
- `MiniBatchDictionaryLearning`

# IPython Notebook:

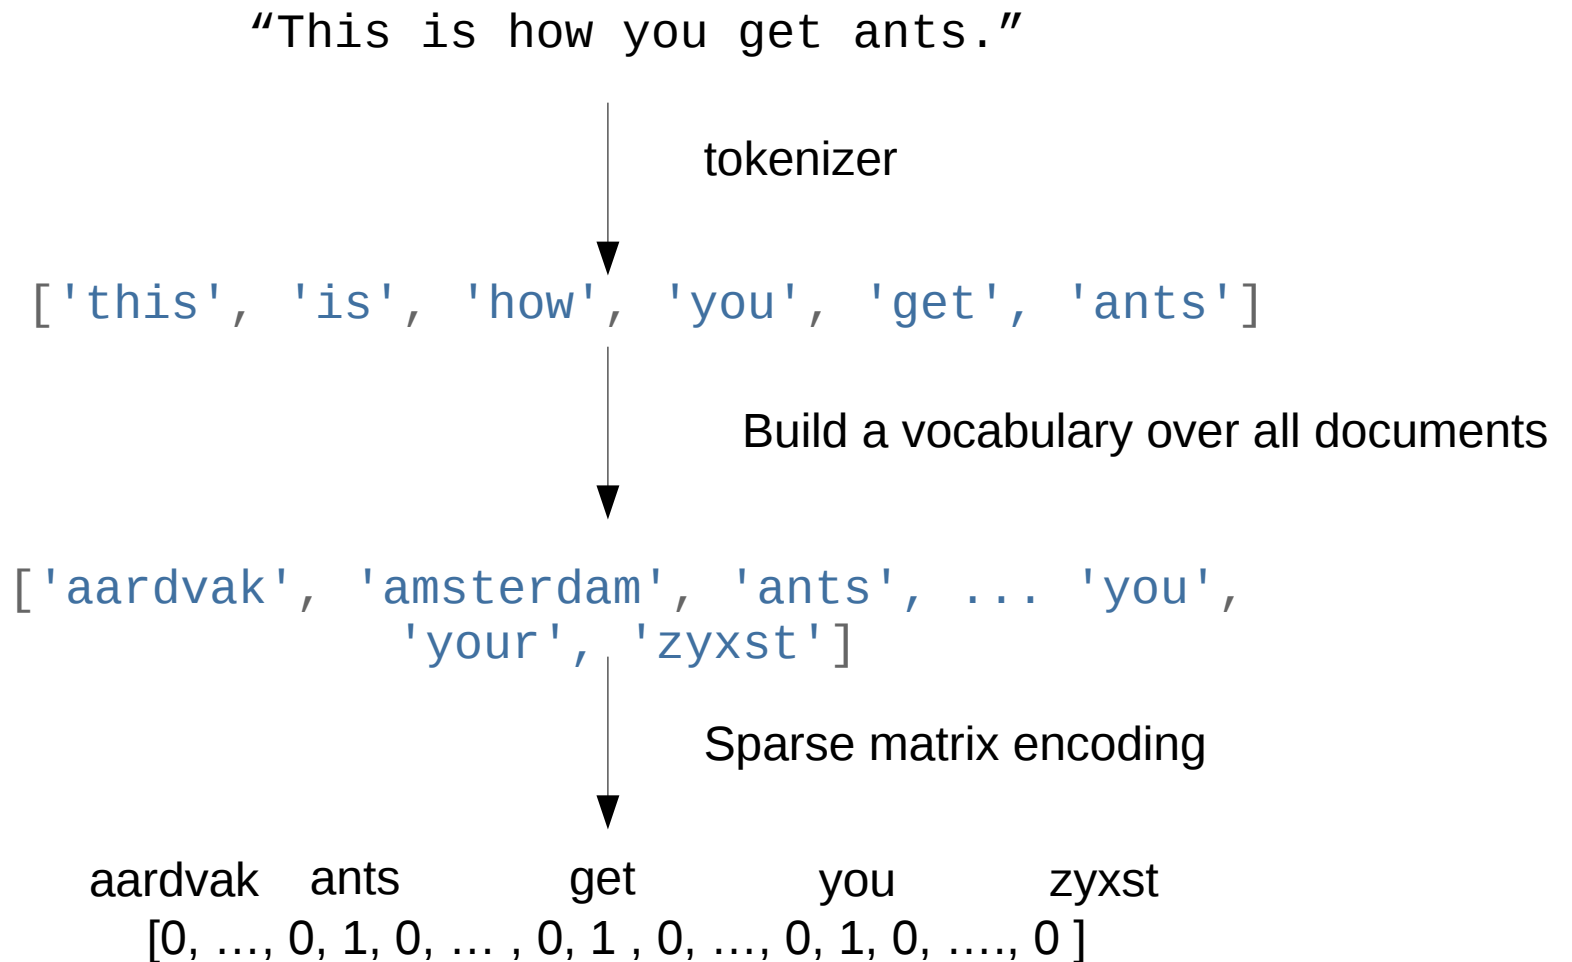
## Part 8 – Out Of Core Learning

# Stateless Transformers

- Normalizer
- HashingVectorizer
- RBFSampler (and other kernel approx)

# Bag Of Word Representations

CountVectorizer / TfidfVectorizer



# Hashing Trick

HashingVectorizer

"This is how you get ants."

tokenizer

['this', 'is', 'how', 'you', 'get', 'ants']

hashing

[hash('this'), hash('is'), hash('how'), hash('you'),  
hash('get'), hash('ants')]  
= [832412, 223788, 366226, 81185, 835749, 173092]

Sparse matrix encoding

aardvak ants get you zyxst  
[0, ..., 0, 1, 0, ..., 0, 1, 0, ..., 0, 1, 0, ..., 0]

# IPython Notebook:

## Part 9 – Out Of Core Learning for Text

# Video Series


## Advanced Machine Learning with scikit-learn

50% Off Coupon Code: AUTHD

# Video Series

## Advanced Machine Learning with scikit-learn

50% Off Coupon Code: AUTHD






# Thank you for your attention.



@t3kcit



@amueller



importamueller@gmail.com



<http://amueller.github.io>